

# Engine Tuning with Msqdev

Jeremiah Mahler

January 21, 2012

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Installation . . . . .	2
1.2	Managing Modifications with Git . . . . .	3
<b>2</b>	<b>Steady State Tuning</b>	<b>3</b>
2.1	convergence of VE tuning . . . . .	3
<b>3</b>	<b>Acceleration Tuning</b>	<b>7</b>
<b>4</b>	<b>Air Fuel Ratio</b>	<b>7</b>
4.1	Predicting Air Fuel Ratios . . . . .	8
4.2	Finding The Air Fuel Ratio For Maximum Power . . . . .	8
<b>A</b>	<b>afr_predict.R</b>	<b>8</b>

# 1 Introduction

The Megasquirt (Bowling & Grippo, 2011b) engine controller allows for the control of nearly every possible parameter of a running engine. But this wide spectrum of parameters also makes it difficult to tune. The typical tuning method (Bowling & Grippo, 2011a) is based on rules of thumb and imprecise judgements of performance. The Msqdev system (Mahler, 2011) was created to provide a methodical method of tuning with provable metrics of performance.

The design of Msqdev consists of various parts. In general small robust programs are preferred over large all encompassing ones. The `msqdev` daemon creates the interface to the ecu. It monitors a directory of files which describe the settings and communicates over the serial port to the Megasquirt ecu when these files are changed or need to be updated. There are utilities such as `msq-ve_tuner` for automatically tuning the fuel mixture and `msq-accel_tuner` for performing acceleration tuning as well as many others. The utilities are small since they only read and write files and it is easy to create new ones. Plotting and analysis of captured data is most often done using R (R Development Core Team, 2011) scripts.

## 1.1 Installation

This project is still in early development and the installation described here allows for updates to be easily performed.

The first step is to retrieve a copy of the current project. The best way is by using Git (Torvalds, 2011) with the command below. If Git is not available the package can be downloaded and unpacked directly. The recommended placement is in the users directory (`~/msqdev`) which is the default if the git clone command is run in the home directory.

```
git clone https://github.com/jmahler/msqdev.git
```

After the Msqdev package has been successfully retrieved and unpacked it is necessary to compile the `msqdev` daemon. This is done by changing to the `msqdev/src/` directory and typing `make`. If all goes correctly the `msqdev` executable will be produced.

In order to access the binaries from any location it is necessary to configure the users `PATH`. The recommended way is to add the `~/bin` to the `PATH` environment variable and then to create symbolic links from inside `bin` directory to those in `msqdev`.

```
jeri@bishop ~$ ls ~/bin | grep msq
msq-accel_tuner
msq-afr_table
msqdev
msq-point_values
msq-rtmonitor
msq-speed_tuner
msq-table_points
msq-ve_tuner
jeri@bishop ~$ grep PATH ~/.bashrc
export PATH="$HOME/bin:$PATH"
jeri@bishop ~$ env | grep PATH
PATH=/home/jeri/bin:/usr/local/bin:/usr/bin:/bin
ri/QuizMaker/bin:/home/jeri/QuizMaker/bin
```

The final step is to verify that `msqdev` works. The following is an example of the steps used on one system.

```
jeri@bishop ~$ mkdir tmp-msqdev
jeri@bishop ~$ cd tmp-msqdev
```

```

jeri@bishop tmp-msqdev$ ls # directory is empty
jeri@bishop tmp-msqdev$ msqdev /dev/ttyUSB0
Mon Jul 4 16:34:17 2011: start
Mon Jul 4 16:34:31 2011: updating files from ecu
^Cjeri@bishop tmp-msqdev$ # quit by Ctrl-C
jeri@bishop tmp-msqdev$ ls # msqdev created files
advanceTable1 afrTable1 pid rtdata veTable1

```

## 1.2 Managing Modifications with Git

One benefit of a file based system is that changes can be managed using a version control system such as Git (Torvalds, 2011). The author's preference is to use one directory for each install. And to use one branch for every tuning session. As an example the first branch could be the 'stock' settings. Another branch could be the settings on a particular date. A benefit of branches is that it is easy to change between them using a simple 'git checkout'.

```

jeri@bishop RedTruck-msqdev$ ls
advanceTable1 afrTable1 analyze pid README.md rtdata veTable1
jeri@bishop RedTruck-msqdev$ git branch
* 2012-10-19
  master
  stock
jeri@bishop RedTruck-msqdev$ git checkout stock
Switched to branch 'stock'
jeri@bishop RedTruck-msqdev$ git branch
  2012-10-19
  master
* stock
jeri@bishop RedTruck-msqdev$

```

## 2 Steady State Tuning

Steady state tuning is done when the engine rpms are stable such as at idle. Its use under heavier load conditions is possible but it is difficult to balance load exactly to the power output.

Steady state tuning is performed by altering variables and then measuring the new steady state. As an example consider varying the fuel mixture at idle to maximize rpm. In a range around the current mixture each setting has a corresponding rpm (see Figure 1). The setting closest to the maximum is the best setting.

The order in which the variables are varied can be helpful in detecting if there is any error caused by acceleration. Suppose, for example that the engine rpms were decelerating. If the variables were varied from one boundary to the other it would have decreasing slope. If instead the variables were varied from the center to one boundary and then back to the center to the other boundary the center point will be separated if there was any acceleration (Figure 2).

### 2.1 convergence of VE tuning

It would be useful to know whether the changes that are being made converge to some stable value. Suppose the goal is to achieve some air fuel ratio and we are in a limited range of the ve map, such as at idle. The changes of each of these points can be plotted over time and if they stabilizing they will reach a constant value.

To display this data it needs to be arranged in a certain way. In this example we will be using R(R Development Core Team, 2011) and Ggplot2(Wickham, 2009).

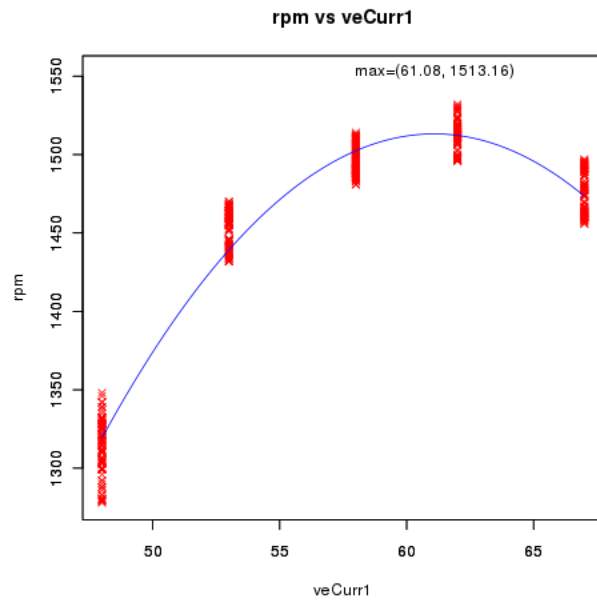


Figure 1: Steady State Tuning by varying veCurr1 vs rpm at idle.




Figure 2: Steady state tuning with a large deceleration error indicated by the gap in the center..

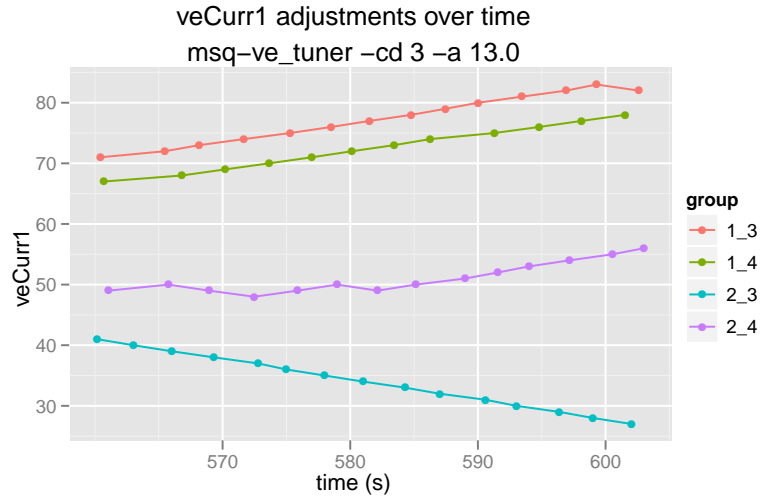


Figure 3: Adjustments of 4 points over time. It can be seen that progress was steady but the time was too short to reach a stable point where it would oscillate.

The data file should consist of the following values:

- time - Time in seconds.
- veCurr1 - The current ve setting.
- group - A unique identifier for a point such as X-Y, where X and Y are the unique row and column.

One possible shortcoming is the fact that points do not change evenly over time. A point that is rarely used will rarely change. Figure 3 and Figure 4 give some examples.

Another problem with auto tuning is that it only covers a subset of the table. How are values for the other points extrapolated? One solution would be to find a least squares fit of the best points found and then use this function to define the remaining points.




Figure 4: Adjustment of many points over time. It is difficult to discern what is going on here and many adjustments look almost random. The point around 3\_ is interesting because it is adjusted way up and then way down. This suggests that it was unable to find a stable point. See Figure 5 for the map at the end of this session. The points are zero offset starting at the bottom left corner at 0\_0.

	RPM															
	[501]	[801]	[1101]	[1401]	[2001]	[2601]	[3101]	[3700]	[4300]	[4900]	[5400]	[6000]	[6500]	[7000]	[7200]	[7500]
[100]	78	85	90	95	99	101	104	109	115	118	117	112	110	108	106	104
[98]	76	84	89	94	98	100	103	108	114	117	116	111	109	107	105	103
[95]	75	83	88	93	97	99	102	107	113	116	115	110	108	106	104	102
[90]	74	82	87	92	96	98	101	106	112	115	114	109	107	105	103	101
[85]	74	81	86	91	95	97	100	105	111	114	113	108	106	104	102	100
[80]	72	77	82	87	90	93	95	100	105	109	108	103	101	99	97	95
[75]	68	70	73	78	81	83	86	90	94	98	97	93	91	89	87	85
[70]	65	66	69	73	76	78	82	86	90	93	92	88	86	84	82	80
[65]	61	62	65	69	72	75	79	82	85	89	88	84	82	80	78	76
[60]	57	59	61	65	69	72	76	78	81	85	85	80	78	76	74	72
[55]	56	59	60	60	62	67	71	73	75	79	77	71	69	67	65	63
[50]	57	60	54	62	61	65	69	71	73	77	75	69	67	65	63	61
[45]	57	63	46	63	60	63	66	69	71	75	73	67	65	63	61	59
[40]	57	56	56	60	60	55	65	67	69	73	71	66	64	62	60	58
[35]	57	56	54	65	61	56	56	60	63	66	65	62	60	58	58	57
[30]	57	56	55	55	59	60	50	55	58	61	61	60	59	59	58	58

Figure 5: The final VE table at the end of the cve tuning session shown in Figure 4.

### 3 Acceleration Tuning

In contrast to Steady State Tuning (Section 2), Acceleration Tuning is done when the engine rpms are increasing or decreasing. This method is more useful than Steady State Tuning because it can be applied to far more situations and does not depend on balanced conditions.

One source of error with acceleration tuning is with the throttle position. A human operator is quite imprecise and can easily allow significant throttle movement without their knowledge. To reduce this error a throttle stop can be used. The throttle stop is a fixed length object such as a piece of tubing which can be placed underneath the gas pedal to limit its travel. Support in the tuning program can be provided to start/stop recording so that a trial is only recording while the pedal is against the stop.




Figure 6: Comparison of two acceleration tuning trials. Each trial is constrained to the same range of rpms. Because the second trial reached the given rpm in a shorter time period it maximized acceleration compared to the first trial.

As can be seen from Figure 6 Acceleration Tuning is good at discerning differences over a range which result in a cumulative difference. It would be very poor at discerning a few marginal adjustments. The values should be varied by considering this characteristic. When adjusting the ve table, for example, multiple values over the range should be adjusted. If the adjustments could be made to produce a constant air fuel ratio over a trial this could be beneficial in finding the ratio that produces the greatest acceleration. Predicting air fuel ratios will be discussed in Section 4.1.

### 4 Air Fuel Ratio

The fuel mixture is crucial in determining how well an engine runs. If it is too rich or too lean the engine will lack power. To minimize emissions the stoichiometric mixture is best. To maximize power a richer

than stoichiometric mixture is best. To maximize fuel economy a leaner than stoichiometric mixture is best. Obviously only one of these can be maximized at a time for a given point on the map. In general the goal of any of these is to achieve a specific fuel mixture for specific operating conditions.

The desired air fuel ratio can be found by incrementally adjusting values while the engine is running until the desired value is achieved. This procedure also results in many data samples of settings and their resulting air fuel ratios. To go beyond merely achieving a single desired air fuel ratio, this data can be used to build a function which predicts what the air fuel ratio will be. This has the added benefit of being able to instantly set the air fuel ratio to any desired value.

## 4.1 Predicting Air Fuel Ratios

The first step in building a function which can predict the air fuel ratio is to record data. This is accomplished in the Msqdev system using `msq-ve_tuner` utility. This utility will attempt to adjust the current ve values to achieve a specific air fuel ratio. The engine should be run through various rpm ranges and various load values. Whether the desired air fuel ratio is achieved is not important, rather data samples of the air fuel ratio for a given setting are what matter.

Once an adequate data sample has been recorded it can be processed using the R script `afr_predict.R` (Appendix A) which performs a least squares fit to find the coefficients which define the linear function. These coefficients can then be used with the Msqdev utility `msq-afr_table` to build tables with any desired air fuel ratio.

There are several sources of error when predicting air fuel ratios. Acceleration enrichment would skew the results because additional fuel is added which is not accounted for. This error can be eliminated by disabling enrichment or by only moving the throttle slowly. Another source of error is with the delay between when settings are executed and when the oxygen sensor reads the value. Wide band Oxygen sensor response times are typically very fast it is estimated that the delay would be no more than a half second. More work is needed to determine the magnitude of these errors and how they can be compensated for.

## 4.2 Finding The Air Fuel Ratio For Maximum Power

Once the function which predicts the air fuel ratio has been found (Section 4.1) the values can be varied to find which one produces the maximum power (maximum acceleration). In the Msqdev system this is accomplished using the utility `msq-accel_tuner` for performing acceleration tuning (Section 3).

A good general starting point for tuning the air fuel ratio is to configure the table so that a constant air fuel ratio is produced during a trial. Then the air fuel ratio can be adjusted during subsequent trials and their performance compared.

## A afr\_predict.R



```

# This script is used to perform a least squares fit of the data
# related to veTable1 to find the function that predicts the
# air fuel ratio.

# To run this script start R and source this file
#
# bash$ R
# > source("afr_predict.R")
# Loading required package: MASS
#
# Call: rlm(formula = afr1 ~ veCurr1 + fuelload + rpm, maxit = 40)
# Residuals:
#      Min       1Q   Median       3Q      Max
# -4.1269 -0.4980 -0.1448  0.5697  3.0852
#
# Coefficients:
#              Value      Std. Error t value
# (Intercept)  17.9873      0.0925   194.3936
# veCurr1      -0.1399      0.0023  -61.2773
# fuelload      0.0525      0.0016   33.8201
# rpm           0.0004      0.0000   25.2164
#
# Residual standard error: 0.7837 on 12986 degrees of freedom
# >

# Once the coefficients are found a table can be built using
# the command below (with different values).
#
# bash$ msq-afr_table -afr 16.5 -a -0.1399 -b 0.0525 -c 0.0004 -d 17.9873

# Choose the file with the recorded data from msq-ve_tuner
#f1 <- "20110626-ve_tuner/msq-ve_tuner-20110626-16:47:48" # ok
f1 <- "20110626-ve_tuner/msq-ve_tuner-20110626-16:59:22" # good
#f1 <- "20110626-ve_tuner/msq-ve_tuner-20110626-17:12:36" # ok

d1 <- read.csv(file=f1, head=TRUE, sep=",")

# Filter the data to remove invalid extreme values.

# Remove invalid extreme air fuel ratios
filt0 <- d1$afr1 > 10 & d1$afr1 < 16

# remove invalid idle and deceleration values
filt0 <- filt0 & d1$tps > 1

# remove near idle rpms
filt0 <- filt0 & d1$rpm > 2600

veCurr1 <- d1$veCurr1[filt0] # x

```

```

afr1    <- d1$afr1[filt0]    # y
fuelload <- d1$fuelload[filt0] # z
rpm     <- d1$rpm[filt0]

# account for delay in afr1
# TODO - how does the filtering (above) affect the delay? Is it correct?
dX <- 0 # ok
#dX <- 10 # ok
#dX <- 20 # ok
#dX <- 40
#dX <- 60 # too far
# shift forward
afr1 <- afr1[(dX + 1):length(afr1)]
# shift back
veCurr1 <- veCurr1[1:(length(veCurr1) - dX)]
fuelload <- fuelload[1:(length(fuelload) - dX)]
rpm <- rpm[1:(length(rpm) - dX)]

#lmfit1 <- lm(formula = afr1 ~ veCurr1 + fuelload + rpm)
require(MASS) # rlm
lmfit1 <- rlm(formula = afr1 ~ veCurr1 + fuelload + rpm, maxit=40)

print(summary(lmfit1))

```

## References

- Bowling, B. & Grippo, A. (2011a). Megamanual. <http://www.megamanual.com>.
- Bowling, B. & Grippo, A. (2011b). Megasquirt. <http://www.bgsflex.com/megasquirt.html>.
- Mahler, J. (2011). Msqdev: Megasquirt development system. <https://github.com/jmahler/msqdev>.
- R Development Core Team (2011). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. ISBN 3-900051-07-0.
- Torvalds, L. (2011). Git. <http://git.scm.com>.
- Wickham, H. (2009). *ggplot2: elegant graphics for data analysis*. Use R! Springer.